

Log4j Sourcetype Investigation - Splunk Enterprise

Date: 2025-08-01 **Environment:** Splunk Enterprise **Index:** main **Sourcetype:** log4j

Objective

To investigate events under the log4j sourcetype in the main index, focusing on identifying the types of events recorded, extracting relevant fields, analyzing success vs failure results, and determining peak failure times for root cause analysis.

Initial Exploration

- Identified log4j and access_combined as top sourcetypes in the main index.
- Chose to begin deep-dive investigation with log4j.
- Searched the log data with a base query:

```
index=main sourcetype=log4j
```
- Observed that logs had entries with terms like result="success" or result="failure".

Search | Splunk 10.0.0 x Search | Splunk 10.0.0 x VirusTotal - API Key - kali x MalShare

127.0.0.1:8000/en-US/app/search/search?q=search index%3Dmain sourcetype%3Dzeek_files_

OffSec Kali Linux Kali Tools Kali Docs Kali Forums Kali NetHunter Exploit-DB Google Hacking DB

splunk>enterprise Apps Administrator 4 Messages Settings Activity Help Find

Search Analytics Datasets Reports Alerts Dashboards Search & Reporting

New Search

Save As Create Table View Close

```
index=main sourcetype=zeek_files_log
| rex field=_raw "(?P<md5>[a-fA-F0-9]{32})"
| stats count by md5
| where count == 1
| sort count asc
```

Time range: All time

839,716 events (before 8/1/25 5:22:45.000 AM) No Event Sampling Job

Events Patterns **Statistics (10,000)** Visualization

Show: 100 Per Page Format Preview: On

md5	count
000034d0a821a040ebe39f5b7b928195	1
00004382922bf280f2f17e7bef1c9bce	1

Search | Splunk 10.0.0 x Search | Splunk 10.0.0 x VirusTotal - API Key - kali x MalShare

127.0.0.1:8000/en-US/app/search/search?q=search index%3Dmain sourcetype%3Dzeek_files_

OffSec Kali Linux Kali Tools Kali Docs Kali Forums Kali NetHunter Exploit-DB Google Hacking DB

splunk>enterprise Apps Administrator 4 Messages Settings Activity Help Find

Search Analytics Datasets Reports Alerts Dashboards Search & Reporting

New Search

Save As Create Table View Close

```
index=main sourcetype=zeek_files_log
| rex field=_raw "(?P<md5>[a-fA-F0-9]{32})"
| rex field=_raw "(?P<mime_type>[a-zA-Z0-9\\-\\/]+)"
| stats count by md5, mime_type
| where count == 1
| sort mime_type, md5
```

Time range: All time

839,716 events (before 8/1/25 5:26:12.000 AM) No Event Sampling Job

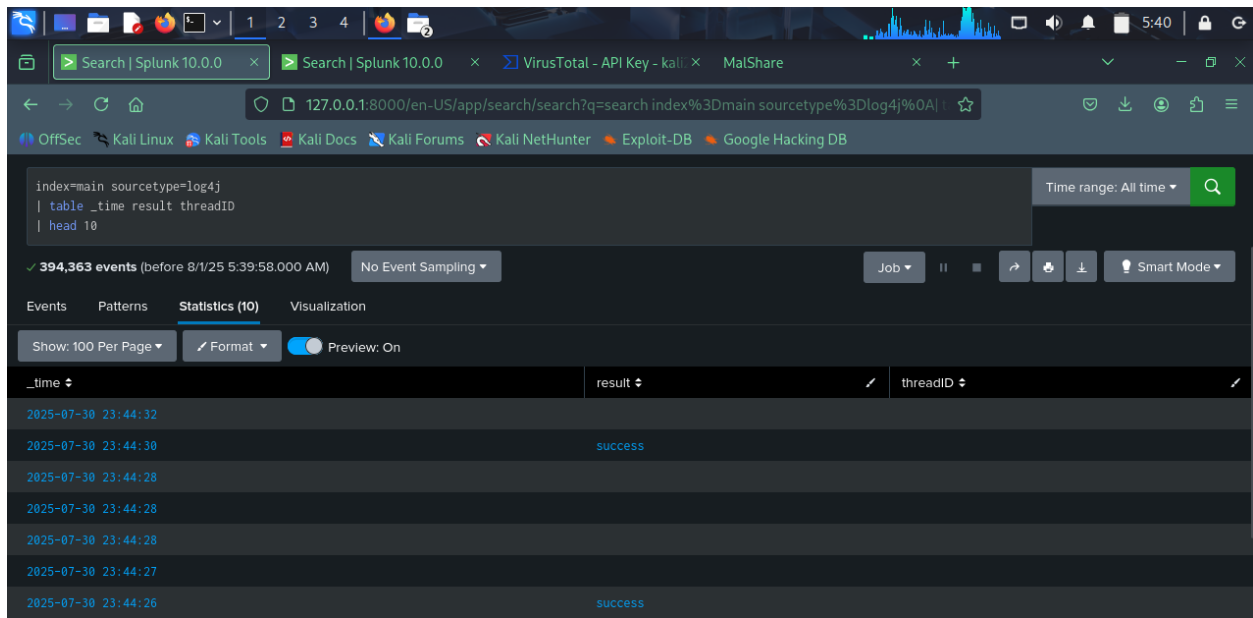
Events Patterns **Statistics (10,000)** Visualization

Show: 100 Per Page Format Preview: On

md5	mime_type	count
d36ef6356fa2aa546f1da2bb003c17b1	1331901001	1

The screenshot displays the Splunk Enterprise web interface. At the top, there's a navigation bar with 'splunk>enterprise' and various user options. Below this is a search bar with the query 'index=main sourcetype=log4j | head 20'. The results section shows '20 events' and a 'Timeline format' view. The first event is a successful response from 'org.webapp.service.shop'.

Time	Event
7/30/25 11:59:59.000 PM	2014-08-04 23:59:59,000 INFO [org.webapp.service.shop] (http--0.0.0.0-8080-47) Response: threadId="10952140719679 5", result="success"



Field Extraction

- Used rex to extract structured fields from the raw log events:


```
| rex "threadId=\"(?<threadId>[^\"]+)\".*?result=\"(?<result>[^\"]+)\| table _time, threadId, result
```
- Initially limited events to 20 to verify field extraction logic.
- Confirmed the appearance of both success and failure values in the result field.

[illegible]

The screenshot shows the Splunk Enterprise interface. The search bar contains the query: `index=main sourcetype=log4j | rex "threadId=\"(?<threadId>[^\"]+)\".*?errorCode=\"(?<errorCode>[^\"]+)\".*?errorMessage=\"(?<errorMessage>[^\"]+)\".*?result=\"(?<result>[^\"]+)\" | table _time threadId result _raw | head 20`. The search results show 7,151 events. The table displays columns: _time, threadId, result, and _raw. The first two rows show failures with error codes 452 and 1738.

_time	threadId	result	_raw
2025-07-30 18:26:53	252361407176807	failure	2014-08-04 18:26:53,000 INFO [org.webapp.service.shop] (http--0.0.0.0-8080-47) Response: threadId="252361407176807", errorCode="452", errorMessage="Out of memory, exiting..." result="failure"
2025-07-30 18:26:43	252361407176798	failure	2014-08-04 18:26:43,000 INFO [org.webapp.service.shop] (http--0.0.0.0-8080-47) Response: threadId="252361407176798", errorCode="1738", errorMessage="Cannot connect to database!" result="failure"

The screenshot shows the Splunk Enterprise interface with the same search query as the previous image. The search results show 7,151 events. The table displays columns: _time, threadId, result, and _raw. The first two rows show failures with error codes 1738 and 1738.

_time	threadId	result	_raw
2025-07-30 18:26:53	252361407176784	failure	2014-08-04 18:26:53,000 INFO [org.webapp.service.shop] (http--0.0.0.0-8080-47) Response: threadId="252361407176784", errorCode="1738", errorMessage="Cannot connect to database!" result="failure"
2025-07-30 18:25:19	252361407176712	failure	2014-08-04 18:25:19,000 INFO [org.webapp.service.shop] (http--0.0.0.0-8080-47) Response: threadId="252361407176712", errorCode="1738", errorMessage="Cannot connect to database!" result="failure"

Filtering Failures

- Modified the query to only include failed events:

```
index=main sourcetype=log4j result="failure"
| rex "threadId=\"(?<threadId>[^\"]+)\".*?errorCode=\"(?<errorCode>[^\"]+)\".*?errorMessage=\"(?<errorMessage>[^\"]+)\".*?result=\"(?<result>[^\"]+)\" | table _time, threadId, errorCode, errorMessage, result
```

- Counted total failures:

- Error Code 1738 (Cannot connect to database): **5108 events**
- Error Code 452 (Out of memory, exiting): **2043 events**

The screenshot shows a Splunk search interface with the following details:

- Search Query:** `| search result="failure"`
`| stats count by errorCode, errorMessage`
`| sort - count`
- Results:** 7151 events (before 8/1/25 6:15:19.000 AM). No Event Sampling.
- Statistics (2):**

errorCode	errorMessage	count
1738	Cannot connect to database!	5108
452	Out of memory, exiting...	2043

⚙️ Field Normalization using coalesce

- Used eval with coalesce to safely extract errorCode in case of field presence inconsistency:
`| eval itercode=coalesce(errorCode, error_code)`
- Ensured that error code statistics aggregated reliably regardless of field name variation.

Search | Splunk 10.0.0 x Search | Splunk 10.0.0 x VirusTotal - API Key - kali x MalShare

127.0.0.1:8000/en-US/app/search/search?q=search index%3Dmain sourcetype%3Dlog4j%0A

OffSec Kali Linux Kali Tools Kali Docs Kali Forums Kali NetHunter Exploit-DB Google Hacking DB

New Search

```
index=main sourcetype=log4j
| rex "threadId=\"(?<threadId>[^\"]+)\"""
| rex "errorCode=\"(?<errorCode>[^\"]+)\"""
| rex "errorMessage=\"(?<errorMessage>[^\"]+)\"""
| eval errorCode=coalesce(errorCode, "none")
| search errorCode="1738"
| table _time threadId errorCode errorMessage
| sort - _time
```

Time range: All time

5,108 events (before 8/1/25 6:21:09.000 AM) No Event Sampling

Job

Events Patterns **Statistics (5,108)** Visualization

Show: 100 Per Page Format Preview: On

_time	threadId	errorCode	errorMessage
2025-07-30 23:59:23	156621407023958	1738	Cannot connect to database!
2025-07-30 23:59:20	109521407196755	1738	Cannot connect to database!
2025-07-30 23:59:14	156621407023950	1738	Cannot connect to database!
2025-07-30 23:59:11	109521407196745	1738	Cannot connect to database!

Search | Splunk 10.0.0 x Search | Splunk 10.0.0 x VirusTotal - API Key - kali x MalShare

127.0.0.1:8000/en-US/app/search/search?q=search index%3Dmain sourcetype%3Dlog4j%0A rex "error:"

OffSec Kali Linux Kali Tools Kali Docs Kali Forums Kali NetHunter Exploit-DB Google Hacking DB

splunk>enterprise Apps Administrator 4 Messages Settings Activity Help Find

Search Analytics Datasets Reports Alerts Dashboards Search & Reporting

New Search

```
index=main sourcetype=log4j
| rex "errorCode=\"(?<errorCode>[^\"]+)\"""
| search errorCode="1738"
| timechart span=1h count as failure_count
```

Time range: All time

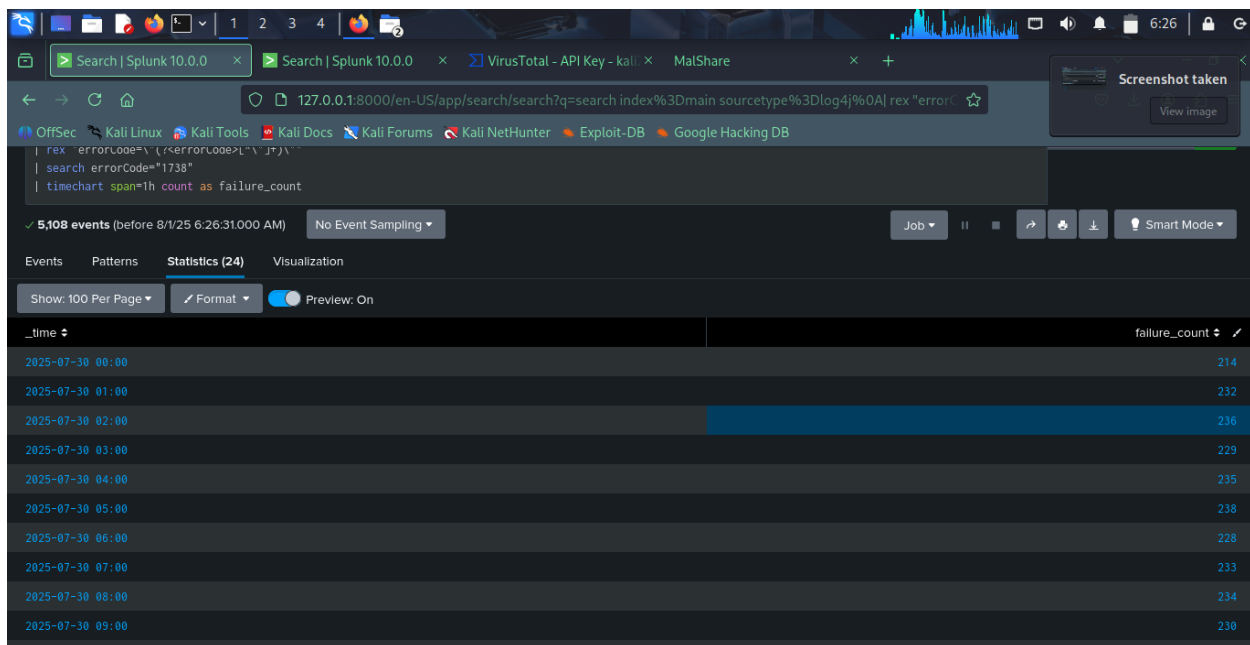
5,108 events (before 8/1/25 6:26:31.000 AM) No Event Sampling

Job

Events Patterns **Statistics (24)** Visualization

Show: 100 Per Page Format Preview: On

_time	failure_count
2025-07-30 00:00	214
2025-07-30 01:00	232
2025-07-30 02:00	236
2025-07-30 03:00	229
2025-07-30 04:00	235



Time-Based Analysis

- Grouped errorCode=1738 failures by hour:

```
index=main sourcetype=log4j errorCode="1738"
| bucket span=1h _time
| stats count as failure_count by _time
```
- Identified **05:00–05:59** as the peak failure hour (highest failure count: **238**).

Deep Dive into Peak Hour

- Queried full events for 05:00–05:59 with errorCode 1738:

```
index=main sourcetype=log4j result="failure" errorCode="1738" earliest="2025-07-30T05:00:00" latest="2025-07-30T05:59:59"
| rex "threadId=\"(?<threadId>[^\"]+)\".*?errorCode=\"(?<errorCode>[^\"]+)\".*?errorMessage=\"(?<errorMessage>[^\"]+)\".*?result=\"(?<result>[^\"]+)\".*?"
| table _time, threadId, errorCode, errorMessage, result
```
- Verified error messages matched expected: Cannot connect to database!
- Analyzed frequency and thread diversity.

Search | Splunk 10.0.0 x Search | Splunk 10.0.0 x VirusTotal - API Key - kali x MalShare

127.0.0.1:8000/en-US/app/search/search?q=search index%3Dlog4j%0A| rex "thres: Time range: All time

```
index=main sourcetype=log4j
| rex "threadId=\"(?<threadId>[^\"]+\""
| rex "errorCode=\"(?<errorCode>[^\"]+\""
| rex "errorMessage=\"(?<errorMessage>[^\"]+\""
| search result="failure" errorCode="1738"
| eval hour=strftime(_time, "%H")
| where hour="05"
| table _time threadId errorCode errorMessage
| sort _time
```

238 events (before 8/1/25 6:30:27:000 AM) No Event Sampling

Events Patterns **Statistics (238)** Visualization

Show: 100 Per Page Format Preview: On

_time	threadId	errorCode	errorMessage
2025-07-30 05:00:01	324431407041996	1738	Cannot connect to database!
2025-07-30 05:00:04	284871407214799	1738	Cannot connect to database!
2025-07-30 05:00:17	284871407214809	1738	Cannot connect to database!
2025-07-30 05:00:55	284871407214849	1738	Cannot connect to database!
2025-07-30 05:01:04	284871407214858	1738	Cannot connect to database!

Summary

A complete investigation was carried out to identify and analyze failure events in the log4j sourcetype. I successfully isolated the most common failure condition (errorCode=1738) and determined its peak hour for occurrence. This prepared me for a more targeted root cause analysis moving forward.